

# Homework

- Lab open M–F 8:00am–4:30pm
- On github.com
  - Join team Everyone
  - Class materials at <http://github.com/Ling583/notes>
  - Read *Bad Data Handbook* Chapter 4: “Bad Data Lurking in Plain Text”
  - Updated python in baddata4.ipynb
- On datacamp.com
  - Finish *Intro to Python for Data Science*
  - Do *Introduction to Shell for Data Science*



#4,935,860

CERTIFICATE NUMBER

# STATEMENT OF ACCOMPLISHMENT

HAS BEEN AWARDED TO

**rmalouf**

FOR SUCCESSFULLY COMPLETING

**Intro to Python for Data Science Course**



**DataCamp**

# Why Unix?

- Big applications vs. toolkit approach
- Easier to document reproducible workflows
- Batch processing
- Open source software
- Widely used across platforms (Linux, cygwin, git-bash, MacOS X)
- File sharing, remote access, security, multiprocessing, etc all pluses

# Why not Unix?

- Cryptic, hard to remember command names
- Vintage 1960's user interface
- Documentation is sparse, confusing, filled with inside jokes, and sometimes wrong
- Open source software
- Multiple incompatible dialects (Linux, cygwin, git-bash, MacOS X)





# Text

- Most basic format is “plain” or “raw” text
- No such thing!
- Computers can only deal with *bits* (ones and zeros), conventionally chunked into *bytes* (8 bits)
- Represent text by giving each letter a numeric *code* between 0 and 255

base 2	01010100	01101000	01101111
base 16	54	68	6F
base 10	84	104	111
	T	h	o

# ASCII

- American Standard Code for Information Interchange
- developed out of 5 bit teletype codes in 1950's and 60's
- unambiguous 7 bit code that used all  $2^7=128$  positions
- included upper and (eventually) lower case letters, numbers, punctuation, control characters (XON/XOFF, CR/LF/FF)
- character codes arranged in collating sequence
- ASCII is (allegedly) suitable only for Latin, Swahili, Hawaiian, and American English, so national variants (ISO-646) were introduced

Dec	Hex	Name	Char	Ctrl-char	Dec	Hex	Char	Dec	Hex	Char	Dec	Hex	Char
0	0	Null	NUL	CTRL-@	32	20	Space	64	40	@	96	60	`
1	1	Start of heading	SOH	CTRL-A	33	21	!	65	41	A	97	61	a
2	2	Start of text	STX	CTRL-B	34	22	"	66	42	B	98	62	b
3	3	End of text	ETX	CTRL-C	35	23	#	67	43	C	99	63	c
4	4	End of xmit	EOT	CTRL-D	36	24	\$	68	44	D	100	64	d
5	5	Enquiry	ENQ	CTRL-E	37	25	%	69	45	E	101	65	e
6	6	Acknowledge	ACK	CTRL-F	38	26	&	70	46	F	102	66	f
7	7	Bell	BEL	CTRL-G	39	27	'	71	47	G	103	67	g
8	8	Backspace	BS	CTRL-H	40	28	(	72	48	H	104	68	h
9	9	Horizontal tab	HT	CTRL-I	41	29	)	73	49	I	105	69	i
10	0A	Line feed	LF	CTRL-J	42	2A	*	74	4A	J	106	6A	j
11	0B	Vertical tab	VT	CTRL-K	43	2B	+	75	4B	K	107	6B	k
12	0C	Form feed	FF	CTRL-L	44	2C	,	76	4C	L	108	6C	l
13	0D	Carriage feed	CR	CTRL-M	45	2D	-	77	4D	M	109	6D	m
14	0E	Shift out	SO	CTRL-N	46	2E	.	78	4E	N	110	6E	n
15	0F	Shift in	SI	CTRL-O	47	2F	/	79	4F	O	111	6F	o
16	10	Data line escape	DLE	CTRL-P	48	30	0	80	50	P	112	70	p
17	11	Device control 1	DC1	CTRL-Q	49	31	1	81	51	Q	113	71	q
18	12	Device control 2	DC2	CTRL-R	50	32	2	82	52	R	114	72	r
19	13	Device control 3	DC3	CTRL-S	51	33	3	83	53	S	115	73	s
20	14	Device control 4	DC4	CTRL-T	52	34	4	84	54	T	116	74	t
21	15	Neg acknowledge	NAK	CTRL-U	53	35	5	85	55	U	117	75	u
22	16	Synchronous idle	SYN	CTRL-V	54	36	6	86	56	V	118	76	v
23	17	End of xmit block	ETB	CTRL-W	55	37	7	87	57	W	119	77	w
24	18	Cancel	CAN	CTRL-X	56	38	8	88	58	X	120	78	x
25	19	End of medium	EM	CTRL-Y	57	39	9	89	59	Y	121	79	y
26	1A	Substitute	SUB	CTRL-Z	58	3A	:	90	5A	Z	122	7A	z
27	1B	Escape	ESC	CTRL-[	59	3B	;	91	5B	[	123	7B	{
28	1C	File separator	FS	CTRL-\	60	3C	<	92	5C	\	124	7C	
29	1D	Group separator	GS	CTRL-]	61	3D	=	93	5D	]	125	7D	}
30	1E	Record separator	RS	CTRL-^	62	3E	>	94	5E	^	126	7E	~
31	1F	Unit separator	US	CTRL-~	63	3F	?	95	5F	_	127	7F	DEL



**Table B-1. ISO Substitution Characters**

ISO	Name	ID	Decimal Character Equivalents											
			35	36	64	91	92	93	94	96	123	124	125	126
6	ASCII	0U	#	\$	@	[	\	]	^	`	{		}	~
4	United Kingdom	1E	£	\$	@	[	\	]	^	`	{		}	-
69	French	1F	£	\$	à	°	ç	§	^	μ	é	ù	è	-
21	German	1G	#	\$	§	Ä	Ö	Ü	^	`	ä	ö	ü	ß
15	Italian	0I	£	\$	§	°	ç	é	^	ù	à	ò	è	ì
11	Swedish for Names	0S	#	□	É	Ä	Ö	Å	Ü	é	ä	ö	å	ü
17	Spanish	2S	£	\$	§	í	Ñ	¿	^	`	°	ñ	ç	-
60	Norwegian version 1	0D	#	\$	@	Æ	Ø	Å	^	`	æ	ø	å	-
2	Int'l. Ref. Version*	2U	#	□	@	[	\	]	^	`	{		}	-
25	French*	0F	£	\$	à	°	ç	§	^	`	é	ù	è	-
	HP German*	0G	£	\$	§	Ä	Ö	Ü	^	`	ä	ö	ü	ß
14	JIS ASCII*	0K	#	\$	@	[	¥	]	^	`	{		}	-
57	Chinese*	2K	#	¥	@	[	\	]	^	`	{		}	-
10	Swedish*	3S	#	□	@	Ä	Ö	Å	^	`	ä	ö	å	-
	HP Spanish*	1S	#	\$	@	í	Ñ	¿	°	`	{	ñ	}	-
85	Spanish*	6S	#	\$	·	í	Ñ	Ç	¿	`	·	ñ	ç	-
16	Portuguese*	4S	#	\$	§	Ä	Ç	Ö	^	`	ä	ç	ö	°
84	Portuguese*	5S	#	\$	·	Ä	Ç	Ö	^	`	ä	ç	ö	-
61	Norwegian version 2*	1D	§	\$	@	Æ	Ø	Å	^	`	æ	ø	å	

# ISO-646

- Not possible to mix languages in a single file with ISO-646
- Loss of punctuation characters a serious drawback:

```
print(word, fdist[word], '\n')
```

becomes in German:

```
print(word, fdist[word], 'Ön')
```

- ISO 646 is a seven bit code, but virtually all computers since the 1960's work with multiples of eight bits

# ISO-8859-1

- ISO-8859-1 is an 8 bit code that extends ASCII to use all  $2^8=256$  positions
- First 7 bits (codes 0–127) are the same as US ASCII
- adds accented and non-Latin characters in codes 128–255
- By convention, lines in text files are separated by LF (10), CR (13) or both
- SPACE (32) and TAB (9) characters used for spacing, indentation, etc

## ISO-8859-1

	x0	x1	x2	x3	x4	x5	x6	x7	x8	x9	xA	xB	xC	xD	xE	xF
0x	<u>NUL</u>	<u>SOH</u>	<u>STX</u>	<u>ETX</u>	<u>EOT</u>	<u>ENQ</u>	<u>ACK</u>	<u>BEL</u>	<u>BS</u>	<u>TAB</u>	<u>LF</u>	<u>VT</u>	<u>FF</u>	<u>CR</u>	<u>SO</u>	<u>SI</u>
1x	<u>DLE</u>	<u>DC1</u>	<u>DC2</u>	<u>DC3</u>	<u>DC4</u>	<u>NAK</u>	<u>SYN</u>	<u>ETB</u>	<u>CAN</u>	<u>EM</u>	<u>SUB</u>	<u>ESC</u>	<u>FS</u>	<u>GS</u>	<u>RS</u>	<u>US</u>
2x	<u>SP</u>	!	"	#	\$	%	&	'	(	)	*	+	,	-	.	/
3x	0	1	2	3	4	5	6	7	8	9	:	;	<	=	>	?
4x	@	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
5x	P	Q	R	S	T	U	V	W	X	Y	Z	[	\	]	^	_
6x	`	a	b	c	d	e	f	g	h	i	j	k	l	m	n	o
7x	p	q	r	s	t	u	v	w	x	y	z	{		}	~	<u>DEL</u>
8x	<u>PAD</u>	<u>HOP</u>	<u>BPH</u>	<u>NBH</u>	<u>IND</u>	<u>NEL</u>	<u>SSA</u>	<u>ESA</u>	<u>HTS</u>	<u>HTJ</u>	<u>VTS</u>	<u>PLD</u>	<u>PLU</u>	<u>RI</u>	<u>SS2</u>	<u>SS3</u>
9x	<u>DCS</u>	<u>PU1</u>	<u>PU2</u>	<u>STS</u>	<u>CCH</u>	<u>MW</u>	<u>SPA</u>	<u>EPA</u>	<u>SQS</u>	<u>SGCI</u>	<u>SCI</u>	<u>CSI</u>	<u>ST</u>	<u>OSC</u>	<u>PM</u>	<u>APC</u>
Ax	<u>NBSP</u>	ı	ç	£	¤	¥	ı	§	"	©	ª	«	¬	<u>SHY</u>	®	™
Bx	°	±	²	³	´	µ	¶	·	,	¹	º	»	¼	½	¾	¿
Cx	À	Á	Â	Ã	Ä	Å	Æ	Ç	È	É	Ê	Ë	Ì	Í	Î	Ï
Dx	Ð	Ñ	Ò	Ó	Ô	Õ	Ö	×	Ø	Ù	Ú	Û	Ü	Ý	Þ	ß
Ex	à	á	â	ã	ä	å	æ	ç	è	é	ê	ë	ì	í	î	ï
Fx	ð	ñ	ò	ó	ô	õ	ö	÷	ø	ù	ú	û	ü	ý	þ	ÿ

# ISO-8859

- ISO-8859-1 (aka “Latin1”) is backwards compatible with US ASCII and covers most Western European languages
- additional code tables added for other languages:

ISO-8859-1 Latin1 (W. Europe)	ISO-8859-8 Hebrew
ISO-8859-2 Latin2 (E. Europe)	ISO-8859-9 Latin5 (Turkish)
ISO-8859-3 Latin3 (Esperanto, Maltese)	ISO-8859-10 Latin6 (Nordic)
ISO-8859-4 Latin4 (N. Europe)	ISO-8859-11 Thai
ISO-8859-5 Cyrillic	ISO-8859-13 Latin7 (Baltic)
ISO-8859-6 Arabic	ISO-8859-14 Latin8 (Celtic)
ISO-8859-7 Greek	ISO-8859-15 Latin9 (W. Europe)

- Accented characters do not fit into blocks or sequence
- Some languages use more than 128 non-ASCII characters

# CJKV scripts

- Character set vs. coded character set vs. encoding
- For all encodings, backwards compatibility with US ASCII is vital
- Some contexts assume eight bit characters, some seven
- Encodings:
  - ISO-2022 represents non-ASCII characters using two seven/eight bit codes
  - Extended Unix Codes (EUC) uses same strategy as ISO-8859

# Problems

- Many, many other encodings, often specific to particular software
- Only languages which share an encoding can be included in the same document
- Not all encodings are easily distinguishable
- Searching and indexing collections of documents with multiple character encodings is impossible
- Not all alphabets use the same collating sequence or text direction
- Stateful encodings
- Mapping between upper and lower case is difficult

# Unicode (ISO-10646)

- Extends ISO-8859-1 to 31 bits ( $2^{31}$ = more than two billion, space for every letter from every writing system ever)
- First 256 positions are the same as ISO-8859-1
- Only 16 bits ( $2^{16}$ =65,534 positions) used at first, with codes for most current world scripts plus numbers, punctuation, diacritics, mathematical symbols, IPA, dingbats, arrows, Braille patterns, Kangxi radicals, etc.
- Now uses 21 bits (more than two million positions), with additional positions used for specialist scripts like hieroglyphics and cuneiform
- Defines a standard name and code for every character



# Unicode (ISO-10646)

- Codes are organized into **blocks** and **scripts**
- Includes composing characters that allow many more characters to be constructed but introduce ambiguities (e.g., Ä could be encoded as the single character U+00C4 or the sequence U+0041 U+0308)
- Extensive character database
  - Case mappings
  - Collation orders
  - Normalization tables
  - Properties: Letter (uppercase, lowercase, titlecase, other), Marks, Number, Punctuation, Symbol, Separator, Other

# Unicode (ISO-10646)

- Offers round trip compatibility with existing national character sets
- Defines standard methods for collation and case folding
- Includes support for multiple text directions (left to right and right to left, but not top to bottom, bottom to top, or boustrophedonic)
- Space reserved for private use characters (Vai, Ethiopic, Klingon)
- Assigns a code position to each character, but doesn't impose an encoding: two simple representations are UCS-2 (aka UTF-16) and UCS-4 (aka UTF-32), just the first 16 or 32 bits of the Unicode character code

# UTF-8

- Unicode solves some of the problems of ISO-8859, but:
  - It doubles the size of characters: *q* is 71 in ISO-8859, U+0071 in UTF-16 and U+00000071 in UTF-32
  - Programs written for 8 bit characters may misinterpret parts of 16 bit characters: e.g., 00 may mean end of string

# UTF-8

- One solution is UTF-8:
  - UTF-8 is backwards compatible with ISO-8859-1, can represent all 32 bits of UCS-4, and preserves the UCS-4 collating sequence.


Bits of code point	First code point	Last code point	Bytes in sequence	Byte 1	Byte 2	Byte 3	Byte 4	Byte 5	Byte 6
7	U+0000	U+007F	1	0xxxxxxx					
11	U+0080	U+07FF	2	110xxxxx	10xxxxxx				
16	U+0800	U+FFFF	3	1110xxxx	10xxxxxx	10xxxxxx			
21	U+10000	U+1FFFFF	4	11110xxx	10xxxxxx	10xxxxxx	10xxxxxx		
26	U+200000	U+3FFFFFFF	5	111110xx	10xxxxxx	10xxxxxx	10xxxxxx	10xxxxxx	
31	U+4000000	U+7FFFFFFF	6	1111110x	10xxxxxx	10xxxxxx	10xxxxxx	10xxxxxx	10xxxxxx

# UTF-8

- UTF-8 is the default encoding for most systems now
- Plain text files usually use LF (`\n`) to mark end of line and SPACE and TAB (`\t`) for spacing
- Python3 (but not Python2!) uses Unicode internally with UTF-8 as the default input and output encoding
- If you are creating the file or if it's well documented, all is well
- Different encodings can be specified on input and output
- Convert Latin1 to EBCDIC

```
for line in open('infile.txt', encoding='latin1'):
    print(line, encoding='cp500')
```

# Mojibake

[illegible]

[a/sa/4a,](#)
[a/2a/4a"](#)
[e-2e5](#)
[ç"en†](#)
[âttæ'â"çµ°](#)
[57](#)
[11](#)

[illegible][illegible]

1 ä, ãl "äZŸä;

[illegible]

1.3 æ-~~t~~a-a, ä/<sup>u</sup>ä, ä/%ä/fä,ä,<sup>u</sup>ä, ä! @ä%ä! ä! æ-~~c~~äl <sup>u</sup>ä,ä/f/ä/@ä/-ä/%

1.4  $\tilde{a}f - \tilde{a}f\tilde{a}, {}^{\circ}\tilde{a}f @ \tilde{a}f \tilde{a} \text{ } @ \tilde{a} - \vee \tilde{a} \tilde{a} \tilde{a} - \tilde{e}^{\circ} \tilde{a}^{\circ} \tilde{a}^{\circ} \tilde{a}^{\circ} \tilde{a}^{\circ} \tilde{a}^{\circ} \tilde{a}^{\circ} @ \tilde{c}^{\circ} \tilde{a} \text{ } \cdot \tilde{a} \text{ } \cdot \tilde{a}, \tilde{a}, \tilde{a}f^{\circ} \tilde{a}f @ \tilde{a}f - \tilde{a}f \text{ }.$

1.5 ééðäçµœë~š0 şă0 @ăf-ăf@ăf-ăf\*

# Mojibake

- Even modern files in Latin scripts run into problems
- Smart quotes and dashes

—	→	â€”
“	→	â€œ
”	→	â€™\x9d

- Accents

schön	→	schÃ¶n
-------	---	--------

- Ligatures

flubberific	≠	flubberific
-------------	---	-------------

# Mojibake

- Some problems are obvious, though may require knowledge of the (human) language
- In Python3
  - chardet module guesses what encoding a file is in  
<https://chardet.readthedocs.io/>
  - ftfy uses heuristics to fix common problems  
<http://ftfy.readthedocs.io/>
  - general Unicode support  
<https://docs.python.org/3/howto/unicode.html>



# Text

- So, “plain” text is text in some encoding (usually UTF-8) with some line break code (usually \n)
- Create or edit using TextEdit, Notepad, etc.
- Read and write with Python’s built in functions (read, print)
- Lingua franca format – convert from whatever to plain text for further processing

# Text

- Other common “plain” formats for tabular data
- CSV (Comma-separated values)

```
Title,Author,ISBN13,Pages
1984,George Orwell,978-0451524935,268
Animal Farm,George Orwell,978-0451526342,144
Brave New World,Aldous Huxley,978-0060929879,288
Fahrenheit 451,Ray Bradbury,978-0345342966,208
Jane Eyre,Charlotte Brontë,978-0142437209,532
Wuthering Heights,Emily Brontë,978-0141439556,416
Agnes Grey,Anne Brontë,978-1593083236,256
Walden,Henry David Thoreau,978-1420922615,156
Walden Two,B. F. Skinner,978-0872207783,301
"Eats, Shoots & Leaves",Lynne Truss,978-1592400874,209
```

- Tab separated files use TAB character (\t) instead of comma
- Create and editing using Excel, etc
- In Python, use csv module or pandas

# Text

- Rich Text Format (RTF)
- Microsoft proprietary formats
  - Word (.doc, .docx)
  - Excel (.xls, .xlsx)
  - Powerpoint (.ppt, .pptx)
- Convert to plain text using “Export” functions
- Convert Word to Text using catdoc
- Read Excel files using pandas

# Text

- Searchable PDFs can be converted to text using pdftotext
- Non-searchable PDFs need to be made searchable first (Adobe Acrobat Pro)
- Ebooks (mobi, epub djvu) via Calibre
- More esoteric task-specific formats
  - Email: rfc822, outlook
  - Geodata: ARC shapefiles
- Web: HTML
- Serialization
  - XML, JSON